

Movie Question Answering System

Using Retrieval-Augmented Generation and Large Language Models

1. Overview

1.1 Project Goal

Build an intelligent question-answering system that can answer natural language questions about a dataset of 10,000 IMDB movies. Your system must handle two fundamentally different types of queries:

1. **Semantic/Conceptual Queries:** Questions requiring understanding of movie content, themes, and context
 - Examples: "What are some alien-related movies?", "Which films explore time travel themes?", "What movies feature strong female protagonists?"
2. **Factual/Statistical Queries:** Questions requiring numerical computations, aggregations, or filtering
 - Examples: "What's the average rating of James Bond movies?", "Which director has the highest-grossing film?", "How many sci-fi movies were released after 2010?"

1.2 Dataset

You will work with an IMDB movie dataset containing ~10,000 movies.

Dataset Link:

<https://drive.google.com/file/d/1MtthDLw4VXobg9EoUf0Nk5uZlh01e0CR/view?usp=sharing>

2. Technical Requirements

2.1 Suggested Technologies

- **LLM:** Qwen, Llama, or similar models.
- **Embeddings:** sentence-transformers

- **Platform:** Google Colab Pro (required for GPU access)
- **Key Libraries:** transformers, torch, pandas, numpy, llama-index, faiss

2.2 Suggested System Architecture

Your system can include the following components:

1. Data Preprocessing Module

- Load and clean the movie dataset
- Create rich text representations for embedding
- Handle missing values appropriately

2. Vector Index Construction

- Generate embeddings for movie descriptions
- Build and persist a vector index using LlamaIndex
- Implement efficient similarity search

3. Query Classification System

- Automatically classify incoming queries as semantic or factual
- Route queries to appropriate processing pipeline

4. Semantic Query Pipeline (RAG)

- Retrieve relevant movie documents using vector similarity
- Generate contextual answers using the LLM
- Provide source attribution (which movies informed the answer)

5. Factual Query Pipeline (Code Generation)

- Generate pandas/SQL code to answer statistical questions
- Execute code safely with appropriate error handling
- Format numerical results into natural language using the LLM

6. Unified Interface

- Single function to answer any question
- Clear logging of query type and processing steps
- User-friendly output formatting

3. Deliverables

3.1 Code Submission

Submit a well-documented Jupyter notebook containing:

Part 1: Setup and Data Loading (10 points)

- Installation of all required packages
- Model loading with proper quantization
- Dataset loading and preprocessing
- Initial data exploration and statistics

Part 2: Vector Index Construction (10 points)

- Document creation from dataframe
- Embedding generation
- Vector index building
- Index persistence and loading

Part 3: Semantic Query Implementation (25 points)

- RAG pipeline implementation
- Query engine configuration
- Retrieval quality testing
- At least 5 example semantic queries with results

Part 4: Factual Query Implementation (35 points)

- Code generation prompts
- Safe code execution framework
- Result formatting
- At least 5 example factual queries with results

Part 5: Query Classification and Integration (20 points)

- Query classification logic
- Unified question-answering interface
- Comprehensive testing with diverse queries
- Error handling and edge cases